JAHAN KAZIMI

(858) 214-8694 | jahankazimi078@ucla.edu | linkedin.com/in/jahan-kazimi | github.com/jahankazimi078 | jahankazimi078.github.io

EDUCATION

University of California, Los Angeles | *B.S., Statistics & Data Science* Sep 2025 - Jun 2027

- GPA: **3.86**/4.0 | Coursework: Probability, Mathematical Statistics, Statistical Programming, Linear Regression

Los Angeles Pierce College Aug 2024 - Jun 2025

- GPA: **4.0**/4.0, Dean's Honor List | Coursework: Linear Algebra, Calculus I/II/III, Statistics, C++

EXPERIENCE

TVLIP | *Data Science & Machine Learning Intern* Jan 2026 - Apr 2026

- Built and evaluated predictive models (LightGBM, logistic regression, gradient boosting) in Python across **1.2M+** records, running **240+** cross-validated experiments that raised performance from **0.71** to **0.86** macro F1 and increased the reliability of automated case prioritization.
- Through exploratory analysis and feature engineering, found recent activity and historical engagement were the strongest predictors of outcomes, while records with limited history drove nearly **one-third** of model errors; rebuilt features around those gaps and improved the weakest segment by **12%**.
- Ran model validation and error analysis (confusion matrices, calibration curves, cohort slices) and built SQL-backed reporting separating high- and low-confidence predictions, focusing manual review on the riskiest cases and cutting unnecessary review volume by an estimated **18%**.

S&F Services | *Data Science & Analytics Intern* Jun 2025 - Sep 2025

- Built end-to-end data pipelines in Python (Pandas, NumPy) and SQL to ingest, clean, and consolidate **500K+** records from multiple sources into query-ready tables, cutting recurring data-prep time **35%** and saving analysts roughly **12 hours** of manual work each month.
- Traced reporting discrepancies to duplicate records and inconsistent category definitions; standardized a normalized MySQL schema (**6** tables, **20+** columns) with foreign-key and ENUM constraints, reducing recurring data-quality issues by **25%**.
- Trained and evaluated multiclass classifiers (LightGBM, logistic regression) with **5**-fold cross-validation and out-of-fold scoring, reaching **0.83** macro F1 and surfacing the customer characteristics most associated with key outcomes.
- Applied regression and hypothesis testing to quantify trends and segment-level differences, delivering interactive dashboards that gave stakeholders direct access to frequently requested metrics instead of ad-hoc requests.

LAPC Center for Academic Success | *Mathematics and Statistics Tutor* Jan 2025 - Jun 2025

- Tutored Calculus I/II/III, Linear Algebra, and Statistics one-on-one and in small groups, in person and online, and led exam-season workshops built around the topics students struggled with most.

PROJECTS

NeuroGolf 2026 | *Kaggle Bronze Medal, top 10%*

- Solved all **400** ARC-style grid-reasoning puzzles by hand-constructing the smallest exact ONNX model for each, building **113** rule-family builders and using AI coding agents to accelerate exploration and submission iteration.
- Engineered equivalence-gated graph optimizers (constant folding, dead-node elimination, dtype demotion, bit-plane packing) with every rewrite proven output-identical; **103** of **400** final models reduced to a single operator.

UCLA ASA DataFest 2026 | *Team Lead*

- Led a **4**-person team analyzing **7.6M+** hospital encounters to quantify disparities in chronic-disease patient journeys, anchored on a **25,668**-patient Type 2 Diabetes cohort and replicated across four additional conditions.
- Engineered a reproducible Python pipeline deriving journey metrics with bootstrap confidence intervals and nested logistic-regression mediation models; found minoritized patients entered care through the ED at **1.9-3.3x** the rate of White patients after adjusting for social determinants and geography, and built dashboards surfacing those gaps.

NYC 311 City Complaints Intelligence System

- Built an end-to-end analytics pipeline (SQLite + yearly Parquet) and a rule-based NLP taxonomy converting free-text complaint descriptions into standardized issue families, subtypes, and resolution outcomes.
- Trained multiclass resolution-time classifiers on **200K+** complaints (HistGradientBoosting, logistic regression) with feature importance, error slices, and borough-level fairness metrics, finding the slowest borough averaged **27%** longer resolution times than the citywide benchmark.

NBA Shot Quality Analytics Platform

- Built an end-to-end shot-quality pipeline analyzing **654,609** field-goal attempts across three seasons, generating per-shot expected points and **983** player-season shooting profiles.
- Trained a LightGBM xPoints model with game-grouped validation and out-of-fold scoring, finding Points Over Expected held a **0.58** year-over-year correlation, and shipped a public interactive dashboard.

SKILLS

- Languages & Tools: Python (Pandas, NumPy, Scikit-learn, LightGBM), R (Tidyverse, ggplot2, Shiny), SQL, MySQL, Streamlit, Excel, Git, GitHub
- Statistics & Modeling: regression, hypothesis testing, A/B testing, cross-validation, model calibration, experimentation, data visualization